

Logistic Regression Classifier

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import statsmodels.api as sm
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
```

Dataset

In [2]:

```
1 dta = sm.datasets.fair.load_pandas().data
2 dta.head(5)
```

Out[2]:

	rate_marriage	age	yrs_married	children	religious	educ	occupation	occupation_husb	
0	3.0	32.0	9.0	3.0	3.0	17.0	2.0	5.0	0
1	3.0	27.0	13.0	3.0	1.0	14.0	3.0	4.0	3.
2	4.0	22.0	2.5	0.0	1.0	16.0	3.0	5.0	1.
3	4.0	37.0	16.5	4.0	3.0	16.0	5.0	5.0	0.
4	5.0	27.0	9.0	1.0	1.0	14.0	3.0	4.0	4.




Checking for outliers using box plot

In [3]:

```
1 dta.drop(['affairs'],axis=1).plot(kind='box',figsize=(16,8))
```

Out[3]:

<matplotlib.axes._subplots.AxesSubplot at 0x187078d18b0>




Relation between Affair and Occupation of Husband

In [4]:

```
1 fig = plt.figure(figsize=(16,8))
2 sns.scatterplot(x=dta['affairs'],y=dta['occupation_husb'],hue=dta['occupation_husb'])
```

Out[4]:

<matplotlib.axes._subplots.AxesSubplot at 0x18765fea4c0>



We can see that the chances of Divorce differs with the occupation hunbands is in.


Relation between Affair and Occupation of Women.

In [5]:

```
1 fig = plt.figure(figsize=(16,8))
2 sns.scatterplot(x=dta['affairs'],y=dta['occupation'],hue=dta['occupation'])
```

Out[5]:

<matplotlib.axes._subplots.AxesSubplot at 0x18707a698e0>



Women having White Collar Job show high chances of Affair.

Information about Dataset

The dataset contains 6366 observations of 9 variables:

rate_marriage: woman's rating of her marriage

(1 = very poor, 5 = very good)

age: woman's age

yrs_married: number of years married

children: number of children

religious: woman's rating of how religious she is

(1 = not religious, 4 =strongly religious)

educ: level of education

(9 = grade school, 12 = high school, 14 = some college, 16 = college graduate,
17 = some graduate school, 20 = advanced degree)

occupation: woman's occupation

(1 = student, 2 = farming/semi- skilled/unskilled, 3 = "white collar", 4 =
teacher/nurse/writer/technician/skilled, 5 = managerial/business, 6 =
professional with advanced degree)

occupation_husb: husband's occupation (same coding as above)

(1 = student, 2 = farming/semi- skilled/unskilled, 3 = "white collar", 4 =
teacher/nurse/writer/technician/skilled, 5 = managerial/business, 6 =
professional with advanced degree)

affairs: time spent in extra-marital affairs

Making a new Column that will contain Affairs in Binary Form

0 - No Affairs

1 - Having Affairs

In [6]:

```
1 dta['Affairs'] = (dta.affairs > 0).astype(int)
2 dta.head(3)
```

Out[6]:

	rate_marriage	age	yrs_married	children	religious	educ	occupation	occupation_husb	
0	3.0	32.0	9.0	3.0	3.0	17.0	2.0	5.0	0
1	3.0	27.0	13.0	3.0	1.0	14.0	3.0	4.0	3.
2	4.0	22.0	2.5	0.0	1.0	16.0	3.0	5.0	1.

Checking For Null Values

In [7]:

```
1 print(dta.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6366 entries, 0 to 6365
Data columns (total 10 columns):
rate_marriage      6366 non-null float64
age                6366 non-null float64
yrs_married        6366 non-null float64
children           6366 non-null float64
religious          6366 non-null float64
educ               6366 non-null float64
occupation         6366 non-null float64
occupation_husb    6366 non-null float64
affairs            6366 non-null float64
Affairs            6366 non-null int32
dtypes: float64(9), int32(1)
memory usage: 472.6 KB
None
```

No Null values are there in any column. Data type is Float, with affairs as int

Creating Design Matrix

We are creating two design matrices. The first is the matrix of endogenous variables(dependent variables ~ 'X'). The second design matrix is of exogenous variables(Independent Variables ~ 'y'). For this we are using dmatrices function from patsy

In [8]:

```
1 from patsy import dmatrices
2 y,X = dmatrices('Affairs ~ rate_marriage+age+yrs_married+children+religious+educ+C(occu
3 X.head(2)
```

Out[8]:

	Intercept	C(occupation) [T.2.0]	C(occupation) [T.3.0]	C(occupation) [T.4.0]	C(occupation) [T.5.0]	C(occupation) [T.6.0]	C(occu
0	1.0	1.0	0.0	0.0	0.0	0.0	
1	1.0	0.0	1.0	0.0	0.0	0.0	

dmatrices returns a dataframe constituting of categories for categorical variables in our case occupation(female occupation) and occupation_husb(Husbands occupations).

It add a contant variable "intercept" to the dataframe.

In [9]:

```
1 y.head(2)
```

Out[9]:

	Affairs
0	1.0
1	1.0

Converting y into a 1 D flattern array.

In [10]:

```
1 y = np.ravel(y)
2 y
```

Out[10]:

```
array([1., 1., 1., ..., 0., 0., 0.])
```

Renaming Columns of categorical variables in X

for simplicity

In [11]:

```
1 X.rename(inplace=True,columns={'C(occupation)[T.2.0]':'occ_2.0','C(occupation)[T.3.0]':  
2 X.head(5)
```

Out[11]:

	Intercept	occ_2.0	occ_3.0	occ_4.0	occ_5.0	occ_6.0	occ_husb_2.0	occ_husb_3.0	occ
0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	1.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
2	1.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	
3	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	
4	1.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	

In [12]:

```
1 print("Info of X:\n")  
2 X.info()
```

Info of X:

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 6366 entries, 0 to 6365  
Data columns (total 17 columns):  
Intercept      6366 non-null float64  
occ_2.0        6366 non-null float64  
occ_3.0        6366 non-null float64  
occ_4.0        6366 non-null float64  
occ_5.0        6366 non-null float64  
occ_6.0        6366 non-null float64  
occ_husb_2.0   6366 non-null float64  
occ_husb_3.0   6366 non-null float64  
occ_husb_4.0   6366 non-null float64  
occ_husb_5.0   6366 non-null float64  
occ_husb_6.0   6366 non-null float64  
rate_marriage  6366 non-null float64  
age            6366 non-null float64  
yrs_married    6366 non-null float64  
children       6366 non-null float64  
religious      6366 non-null float64  
educ           6366 non-null float64  
dtypes: float64(17)  
memory usage: 895.2 KB
```

Everything looks good. Now let's check whether any values are null values

In [13]:

```
1 X.isna().sum()
```

Out[13]:

```
Intercept      0
occ_2.0         0
occ_3.0         0
occ_4.0         0
occ_5.0         0
occ_6.0         0
occ_husb_2.0    0
occ_husb_3.0    0
occ_husb_4.0    0
occ_husb_5.0    0
occ_husb_6.0    0
rate_marriage  0
age            0
yrs_married    0
children       0
religious      0
educ           0
dtype: int64
```

No Null values are there so Moving forward with our next step.

Checking Correlation

Correlation

Relation of a variable with other variables.

In [14]:

```
1 X.corr()
```

Out[14]:

	Intercept	occ_2.0	occ_3.0	occ_4.0	occ_5.0	occ_6.0	occ_husb_2.0	occ_husb_3.0	occ_husb_4.0	occ_husb_5.0	occ_husb_6.0	rate_marriage	age	yrs_married	children	religious	educ
Intercept	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
occ_2.0	NaN	1.000000	-0.348075	-0.251243	-0.143237	-0.052128	0.183782	-0.020904	-0.009786	-0.093292	-0.059107	-0.019697	-0.034223	0.004668	0.081182	-0.013129	-0.217719
occ_3.0	NaN	-0.348075	1.000000	-0.560645	-0.319631	-0.116322	-0.000638	0.090043	0.011248	0.003021	-0.101673	-0.053082	-0.066371	-0.021261	-0.063298	-0.034986	-0.335615
occ_4.0	NaN	-0.251243	-0.560645	1.000000	-0.230712	-0.083962	-0.083123	-0.043159	0.037341	-0.001946	0.085766	0.068882	0.040982	-0.026816	-0.003235	0.043996	0.477505
occ_5.0	NaN	-0.143237	-0.319631	-0.230712	1.000000	-0.047868	-0.053426	-0.044053	-0.039932	0.114903	0.006016	-0.002109	0.079533	0.076820	0.033274	0.004260	-0.022121
occ_6.0	NaN	-0.052128	-0.116322	-0.083962	-0.047868	1.000000	-0.046140	-0.029028	-0.043541	-0.030926	0.218824	0.008878	0.030676	-0.004912	-0.026830	0.011784	0.226920
occ_husb_2.0	NaN	0.183782	-0.000638	-0.083123	-0.053426	-0.046140	1.000000	-0.146849	-0.347951	-0.316693	-0.153248	-0.038992	-0.057368	-0.033451	0.001190	0.009990	-0.160756
occ_husb_3.0	NaN	-0.020904	0.090043	-0.043159	-0.044053	-0.029028	-0.146849	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
occ_husb_4.0	NaN	-0.009786	0.011248	0.037341	-0.039932	-0.043541	-0.347951	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
occ_husb_5.0	NaN	-0.093292	0.003021	-0.001946	0.114903	-0.030926	-0.316693	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
occ_husb_6.0	NaN	-0.059107	-0.101673	0.085766	0.006016	0.218824	-0.153248	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
rate_marriage	NaN	-0.019697	-0.053082	0.068882	-0.002109	0.008878	-0.038992	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
age	NaN	-0.034223	-0.066371	0.040982	0.079533	0.030676	-0.057368	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000
yrs_married	NaN	0.004668	-0.021261	-0.026816	0.076820	-0.004912	-0.033451	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000
children	NaN	0.081182	-0.063298	-0.003235	0.033274	-0.026830	0.001190	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000	0.000000
religious	NaN	-0.013129	-0.034986	0.043996	0.004260	0.011784	0.009990	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000	0.000000
educ	NaN	-0.217719	-0.335615	0.477505	-0.022121	0.226920	-0.160756	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.000000

VIF Score

VIF Score > 5 Indicates sever correlation.

In [15]:

```
1 # VIF Score
2 from statsmodels.stats.outliers_influence import variance_inflation_factor as vif
3 from sklearn.preprocessing import StandardScaler as StdS
4 scaler = StdS()
5 X_data = scaler.fit_transform(X=X)
6 vif_score = [vif(X_data,i) for i in range(X_data.shape[1])]
7 for i in range(len(X.columns)):
8     print(X.columns[i],":",vif_score[i])
9 dataset1 = X
```

```
Intercept : nan
occ_2.0 : 19.340780081786384
occ_3.0 : 39.33561800583917
occ_4.0 : 32.93190959613787
occ_5.0 : 17.05716507563386
occ_6.0 : 3.697958521454455
occ_husb_2.0 : 5.5662915598555065
occ_husb_3.0 : 2.9910696353751987
occ_husb_4.0 : 6.930281429175677
occ_husb_5.0 : 6.577077423174822
occ_husb_6.0 : 3.1852660411536395
rate_marriage : 1.0387458559035507
age : 5.477889737628554
yrs_married : 7.16961101396522
children : 2.5856914303595504
religious : 1.0375560897654312
educ : 1.6357895734289674
```

```
C:\Users\adity\anaconda3\lib\site-packages\statsmodels\regression\linear_model.py:1717: RuntimeWarning: invalid value encountered in double_scalars
    return 1 - self.ssr/self.uncentered_tss
```

In [16]:

```
1 for i in range(len(X.columns)):
2     print(X.columns[i],":",vif_score[i])
```

```
Intercept : nan
occ_2.0 : 19.340780081786384
occ_3.0 : 39.33561800583917
occ_4.0 : 32.93190959613787
occ_5.0 : 17.05716507563386
occ_6.0 : 3.697958521454455
occ_husb_2.0 : 5.5662915598555065
occ_husb_3.0 : 2.9910696353751987
occ_husb_4.0 : 6.930281429175677
occ_husb_5.0 : 6.577077423174822
occ_husb_6.0 : 3.1852660411536395
rate_marriage : 1.0387458559035507
age : 5.477889737628554
yrs_married : 7.16961101396522
children : 2.5856914303595504
religious : 1.0375560897654312
educ : 1.6357895734289674
```

Now from this VIF Score it is clear that there are features having very high correlation values.

occ_2.0,occ_3.0,occ_4.0 and occ_5.0. In occ_husb - occ_husb_2.0, occ_husb_4.0, occ_husb_5.0 are also having high correlation.


Plotting heat map with correlation value with X.

In [17]:

```
1 fig = plt.figure(figsize=(16,8))
2 sns.heatmap(data=X.corr(),cmap="YlGnBu", annot=True)
```

Out[17]:

<matplotlib.axes._subplots.AxesSubplot at 0x1870a279280>



As we can observe that in occ_2.0, occ_3.0, occ_4.0 and occ_5.0 .

occ_2.0 with {occ_2.0, occ_3.0, occ_4.0 and occ_5.0} is {1.000000 -0.348075 -0.251243 -0.143237} respectively.

occ_3.0 with {occ_2.0, occ_3.0, occ_4.0 and occ_5.0} is {-0.348075 1.000000 -0.560645 -0.319631} respectively.

occ_4.0 with {occ_2.0, occ_3.0, occ_4.0 and occ_5.0} is {-0.251243 -0.560645 1.000000 -0.230712} respectively.

occ_5.0 with {occ_2.0, occ_3.0, occ_4.0 and occ_5.0} is {-0.143237 -0.319631 -0.230712 1.000000} respectively.

The feature occ_3.0 is having high correlation in respect to all other features. So removing occ_3.0 and let's check the effect on VIF Score.

In [18]:

```
1 X_data = scaler.fit_transform(X=X.drop('occ_3.0',axis=1))
2 vif_score = [vif(X_data,i) for i in range(X_data.shape[1])]
3 for i in range(len(X.drop('occ_3.0',axis=1).columns)):
4     print(X.drop('occ_3.0',axis=1).columns[i],":",vif_score[i])
```

```
Intercept : nan
occ_2.0 : 1.177577731855246
occ_4.0 : 1.58134483216193
occ_5.0 : 1.1518453049129829
occ_6.0 : 1.166552590194721
occ_husb_2.0 : 5.530740700962974
occ_husb_3.0 : 2.9761967717120825
occ_husb_4.0 : 6.8868032112916735
occ_husb_5.0 : 6.533261649523526
occ_husb_6.0 : 3.1752920128006568
rate_marriage : 1.038534992602799
age : 5.474644652197843
yrs_married : 7.169399081318488
children : 2.5846493806759088
religious : 1.037500793190026
educ : 1.6292132746346843
```


```
C:\Users\adity\anaconda3\lib\site-packages\statsmodels\regression\linear_model.py:1717: RuntimeWarning: invalid value encountered in double_scalars
    return 1 - self.ssr/self.uncentered_tss
```

In [19]:

```
1 fig = plt.figure(figsize=(16,8))
2 sns.heatmap(data=X.drop('occ 3.0',axis=1).corr(),cmap="YlGnBu", annot=True)
```

Out[19]:

```
<matplotlib.axes. subplots.AxesSubplot at 0x1870b800070>
```



In [20]:

```
1 X.head(2)
```

Out[20]:

	Intercept	occ_2.0	occ_3.0	occ_4.0	occ_5.0	occ_6.0	occ_husb_2.0	occ_husb_3.0	occ_hu
0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	
1	1.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	

In [21]:

```
1 from sklearn.model_selection import train_test_split as tts
2 x_train,x_test,y_train,y_test = tts(X,y,test_size=0.20,random_state=225)
```

Using Stochastic Gradient Descent

In [22]:

```
1 from sklearn.linear_model import SGDClassifier
2 model1 = SGDClassifier()
3 model1.fit(x_train,y_train)
```

Out[22]:

```
SGDClassifier(alpha=0.0001, average=False, class_weight=None,
              early_stopping=False, epsilon=0.1, eta0=0.0, fit_intercept=True,
              l1_ratio=0.15, learning_rate='optimal', loss='hinge',
              max_iter=1000, n_iter_no_change=5, n_jobs=None, penalty='l2',
              power_t=0.5, random_state=None, shuffle=True, tol=0.001,
              validation_fraction=0.1, verbose=0, warm_start=False)
```

Testing Score

In [23]:

```
1 model1.score(x_train,y_train)
```

Out[23]:

0.43676355066771405

Traning Score

In [24]:

```
1 model1.score(x_test,y_test)
```

Out[24]:

0.4340659340659341

Using LogisticRegression Classifier

In [25]:

```
1 from sklearn.linear_model import LogisticRegression
2 model2 = LogisticRegression(solver='liblinear')
3 model2.fit(x_train,y_train)
```

Out[25]:

```
LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                    intercept_scaling=1, l1_ratio=None, max_iter=100,
                    multi_class='auto', n_jobs=None, penalty='l2',
                    random_state=None, solver='liblinear', tol=0.0001, verbos
e=0,
                    warm_start=False)
```

Traning Score

In [26]:

```
1 model2.score(x_train,y_train)*100
```

Out[26]:

73.19324430479183

Testing Score

In [27]:

```
1 model2.score(x_test,y_test)*100
```

Out[27]:

69.78021978021978

Evaluation of Classification Model

1. Accuracy
2. Recall
3. Precision
4. F1 Score
5. Specifity
6. AUC (Area Under Curve)
7. RUC(Receiver Operator Characteristic)

Cross Validation

In [28]:

```
1 # Performing cross validation for model1 and Model 2
2
3 from sklearn.model_selection import cross_val_predict
4 model1_y_predict = cross_val_predict(model1,x_train,y_train,cv=3)
5 model2_y_predict = cross_val_predict(model2,x_train,y_train,cv=3)
```

Confusion Matrix

--	Actual	Values
Predicted	True Positive	False Positive
Values	False Negative	True Negative

For Model 1

In [29]:

```
1 from sklearn.metrics import confusion_matrix,recall_score, precision_score, f1_score, a
2 conf_matrix_model1 = confusion_matrix(y_train,model1_y_predict)
3 conf_matrix_model1
```

Out[29]:

```
array([[3021,  427],
       [1088,  556]], dtype=int64)
```

For Model 2

In [30]:

```
1 conf_matrix_model2 = confusion_matrix(y_train,model2_y_predict)
2 conf_matrix_model2
```

Out[30]:

```
array([[3099,  349],
       [1030,  614]], dtype=int64)
```


In [31]:

```
1 # Accuracy Recall Precision F1 Score Specifity calculator.
2 """
3 def calulate(c_matrix):
4     tp = c_matrix[0][0]
5     fp = c_matrix[0][1]
6     fn = c_matrix[1][0]
7     tn = c_matrix[1][1]
8     acc = (tp+tn)/(tp+fp+fn+tn)
9     rell = tp/(tp+fn)
10    pre = tp/(tp+fp)
11    f1 = 2*(pre*rell)/(pre+rell)
12    spe = tn/(tn+fp)
13    return acc,rell,pre,f1,spe
14 """
15
16 def get_values(model):
17     recall = recall_score(y_train,model)
18     precision = precision_score(y_train,model)
19     f1 = f1_score(y_train,model)
20     accuracy = accuracy_score(y_train,model)
21     return print(f'recall: {recall}\nprecision: {precision}\nf1: {f1}\naccuracy: {accuracy}\nspecificity: {spe}')
22 """
```

```
"""m1_accuracy,m1_recall,m1_precision,m1_f1_score,m1_specificity = calulate(conf_matrix_model1)
print(f'Confusion Matrix:\n{conf_matrix_model1}\nm1_accuracy: {m1_accuracy}\nm1_recall:
{m1_recall}\nm1_precision: {m1_precision}\nm1_f1_score: {m1_f1_score}\nm1_specificity: {m1_specificity}')
```

```
"""m2_accuracy,m2_recall,m2_precision,m2_f1_score,m2_specificity = calulate(conf_matrix_model2)
print(f'Confusion Matrix:{conf_matrix_model2}\nm2_accuracy: {m2_accuracy}\nm2_recall:
{m2_recall}\nm2_precision: {m2_precision}\nm2_f1_score: {m2_f1_score}\nm2_specificity: {m2_specificity}')
```

In [32]:

```
1 print("model 1")
2 get_values(model1_y_predict)
3 print("\nModel 2")
4 get_values(model2_y_predict)
```

```
model 1
recall: 0.3381995133819951
precision: 0.5656154628687691
f1: 0.4232965359725923
accuracy: 0.7024744697564808
```

```
Model 2
recall: 0.3734793187347932
precision: 0.6375908618899273
f1: 0.47103950901419256
accuracy: 0.7291830322073841
```

Now each of these values signifies. Let's find out

Accuracy

Total class values accurately classified by the model. Model 1 Accuracy is about 53 % Model 2 Accuracy is about 72 %

Model 2 is more accurate but in the case of classification only model accuracy is not the criteria of getting an actual accurate model.

Moving forward let's see recall and precision of the model.

Recall

In [33]:

```
1 print("model 1")
2 get_values(model1_y_predict)
3 print("Model 2")
4 get_values(model2_y_predict)
```

```
model 1
recall: 0.3381995133819951
precision: 0.5656154628687691
f1: 0.4232965359725923
accuracy: 0.7024744697564808
Model 2
recall: 0.3734793187347932
precision: 0.6375908618899273
f1: 0.47103950901419256
accuracy: 0.7291830322073841
```

Recall is defined as the ability of ML Model to correctly predict positive out of all positive results. Taking the recall into consideration we can choose Model 1 .

Precision

It measure of amongst all the positive predictions how many of them were were actually positive.

Precision of Model 1 is 37% and of Model 2 is 63% . So if we want that our model must accurately predict the outcome we should consider high precision. taking this into consideration we should go with Model 2.


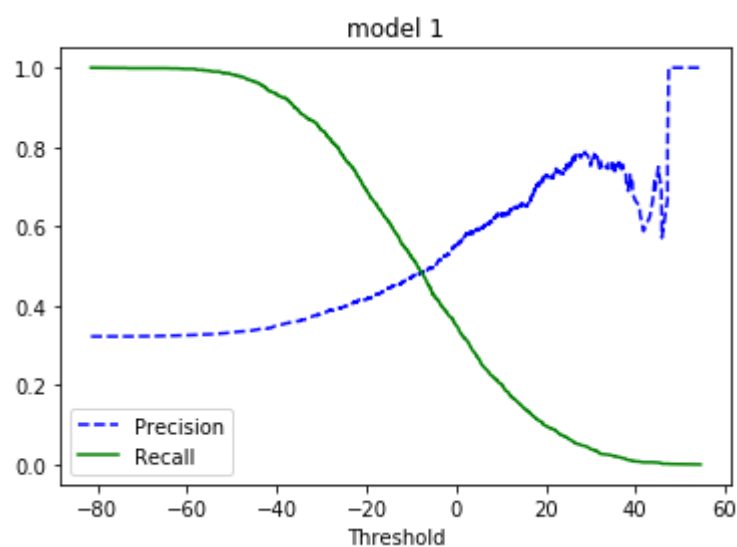
Ploting Precision Recall Curve

In [34]:

```
1 def plot_precision_recall_vs_threshold(precisions, recalls, thresholds,model):
2     plt.title(model)
3     plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
4     plt.plot(thresholds, recalls[:-1], "g-", label="Recall")
5     plt.xlabel('Threshold')
6     plt.legend()
7
```

In [35]:

```
1 from sklearn.metrics import precision_recall_curve as prc
2 y_score_m1 = cross_val_predict(model1,x_train,y_train,cv=3,method='decision_function')
3 y_score_m2 = cross_val_predict(model2,x_train,y_train,cv=3,method='decision_function')
4 pre_m1, rec_m1 , threshol_m1 = prc(y_train,y_score_m1)
5 pre_m2, rec_m2 , threshol_m2 = prc(y_train,y_score_m2)
6
7
8 plot_precision_recall_vs_threshold(pre_m1, rec_m1 , threshol_m1,"model 1")
9 plt.show()
10
11
12 plot_precision_recall_vs_threshold(pre_m2, rec_m2 , threshol_m2,"Model 2")
13 plt.show()
14
```



We can observe with change in threshold we can change the precision and recall of both of our models. But for the time being let's go with the default threshold score.

F1 Score

It's like a trade off factor . classifier having high precision and recall value will get high f1 score. Considering this we should go with Model 1.

ROC

ROC curve is graph of True Positive rate Vs false Positive rate True Positive Rate is Recall/Sensitivity while False Positive rate is out of total false prediction how many were actually false.

AUC

Aea under the curve, According to it the model is best which encloses maximum area.

In [36]:

```
1 from sklearn.metrics import roc_auc_score
2
3 roc_auc_m1 = roc_auc_score(y_train,y_score_m1)
4 roc_auc_m2 = roc_auc_score(y_train,y_score_m2)
5
6 print(f'roc_auc_m1: {roc_auc_m1}\nroc_auc_m2: {roc_auc_m2}')
```



roc_auc_m1: 0.6739852539784691

roc_auc_m2: 0.7477612290491755

Ploting roc_auc

In [37]:

```
1 from sklearn.metrics import roc_curve
2 fpr_m1, tpr_m1, thresholds_m1 = roc_curve(y_train, y_score_m1)
3 fpr_m2, tpr_m2, thresholds_m2 = roc_curve(y_train, y_score_m2)
4
5
6 def plot_roc_curve(fpr, tpr, model, label=None):
7     plt.title(model)
8     plt.plot(fpr, tpr, linewidth=2, label=label)
9     plt.plot([0, 1], [0, 1], 'k--') # dashed diagonal
10    plt.ylabel('Recall')
11    plt.xlabel('FPR')
12    plt.grid()
13    plt.show()
14
15 plot_roc_curve(fpr_m1, tpr_m1, model="Model 1")
16 plot_roc_curve(fpr_m2, tpr_m2, model="Model 2")
17
18
```



According to this plot the best model is the model which is far away from the dotted line. So in this case both our model is good we can use any of the two based on the precision or recall what we want more.

Taking this into consideration Using Model 2 for Application Purpose because we want high precision at 37 % recall

In [38]:

```
1 import pickle
2 file_name = 'Model.pickle'
3 pickle.dump(model2,open(file_name,'wb'))
```

In [40]:

```
1 model = pickle.load(open(file_name,'rb'))
2 #model.predict([a])
```

Testing Score

In [41]:

```
1 model1.score(x_test,y_test)
```

Out[41]:

0.4340659340659341

In [42]:

```
1 model2.score(x_test,y_test)
```

Out[42]:

0.6978021978021978

EOF

In []:

```
1
```